

pp reference manual

a content-addressed, capability-scoped Lisp

Contents

1	Introduction	3
1.1	How to read this manual	4
2	The language in brief	4
2.1	Values	4
2.2	Bindings	4
2.3	Functions	5
2.4	Types	5
2.5	Laziness	5
2.6	Effects and handlers	6
2.7	Config	6
3	Nodes and content-addressing	7
3.1	The node key	7
3.2	Caching across runs	8
3.3	Pure compute needs no capability	9
4	The store and verifying traces	9
4.1	The store	9
4.2	A hit is a verified trace, not a timestamp	10
4.3	pp why	11
4.4	Auditing the cache	12
5	Building with pp	12
5.1	A compilation unit is a node	12
5.2	Depfiles refine the trace	12
5.3	Blobs and materialization	13
5.4	The incremental story	13
6	Capabilities and secrets	15
6.1	Capabilities enter only at the root	15
6.2	User code narrows, never widens	16
6.3	Authority gates the cache, transitively	17
6.4	Secrets: sealed cells	17
7	Modules and islands	19
7.1	Modules	19
7.2	Loading from files	19
7.3	Islands	20
8	Domains and the reconciler	23
8.1	A domain is observe / diff / apply	23
8.2	Reconciling a filesystem	23
8.3	Watching, and other domains	24
8.4	Fenced effects	25
9	Distribution	25
9.1	Identity is location-independent	25
9.2	What is real, and what is stubbed	28

9.3	Host-qualified identity	28
9.4	Trace merge, audit, and growth	28
10	The two back ends	28
10.1	Running each engine	29
10.2	--diff: check them against each other	29
10.3	The differential fuzzer	30
11	Language reference	31
11.1	Value types	31
11.2	Arithmetic and comparison	32
11.3	Booleans and conditionals	32
11.4	Bindings	33
11.5	Functions	34
11.6	Sequencing	34
11.7	Type annotations	35
11.8	Quotation	35
11.9	Laziness	36
11.10	Effects and config	36
11.11	Macros	37
11.12	Modules	37
11.13	Lists	38
11.14	Maps	39
11.15	Strings	40
11.16	Type predicates	41
12	Command-line reference	41
12.1	Running programs	42
12.2	Back ends	42
12.3	Capabilities and grants	42
12.4	The store and auditing	43
12.5	Reconcile, watch, and supervise	43
12.6	Islands	43
12.7	Scheduling and distribution	44
12.7.1	Internal transport (not for hand use)	44
12.8	Observation pinning	45
13	Builtins and the standard library	45
13.1	Core builtins	45
13.1.1	Arithmetic and comparison	45
13.1.2	Lists and pairs	46
13.1.3	Vectors, maps, and sets	46
13.1.4	Type predicates	46
13.1.5	Strings and numbers	47
13.1.6	I/O and control	47
13.1.7	Filesystem observation	47
13.1.8	Hashing and the content store	47
13.1.9	Capabilities	47
13.1.10	Effects, domains, and metaprogramming	48
13.2	Lists	48
13.3	Maps	49
13.4	Strings	49
13.5	Domains	49

13.5.1	Filesystem policy internals (domain-fs.pp)	49
13.5.2	Process policy internals (domain-proc.pp)	50

1 Introduction

pp is a content-addressed, capability-scoped Lisp. It is also an experiment: that the language you write a program in can be the same language you build, package, and deploy it with. A build system, a package manager, and a cluster orchestrator are all the same thing underneath. Each is a graph of computations with caching and side effects. pp makes that graph the language's own evaluation model. So you get incremental builds, deduplication, and reproducible deployment as consequences of how values work, not as separate tools added on top.

Two ideas carry the whole design. First, every value has a content hash. So two computations with the same code and inputs are the same computation. Second, side effects require capabilities: unforgeable tokens, granted only at the command line, that code must hold to touch the world. Everything else in this manual follows from those two ideas.

The fastest way to learn pp is to read it running. Every example here is a real program that the manual's own build ran with pp. The output shown is what it actually printed. Here is the smallest example: values, a conditional, and arithmetic.

hello.pp

```
# Every value has a content hash; arithmetic and comparison are variadic.
print(+(1, 2, 3))
print(if 5 > 0 { "positive" } else { "negative" })
print(6 * 7)
```

```
$ pp hello.pp
```

```
6
"positive"
42
```

The output appears below the command because the manual ran that file. If an example ever stopped producing the output shown, the build for this page would have failed.

That extends to mistakes. pp checks optional type annotations when a function's body runs. Feed one the wrong type and it says so, naming the offending value:

type-error.pp

```
# Type annotations are optional, and checked when the body runs. Passing a
# string where an int is annotated is caught – the manual shows the real error.
def inc(n: int): int { n + 1 }
inc("oops")
```

```
$ pp type-error.pp
```

```
pp: error: type mismatch: expected int, got "oops" at type-error.pp:3
```

You will see errors this way throughout the manual: provoked and shown, not described. So you know exactly what pp does when something is wrong, not just when everything is right.

1.1 How to read this manual

The chapters move from the surface inward. Values, bindings, and functions come first, then laziness, then the effect-and-capability system that governs side effects. From there the manual turns to what makes pp a build substrate: content-addressed nodes and the store, modules and islands, domains and the reconciler, and finally distribution across machines. We introduce each construct, give a sentence or two of motivation, then show it by running it.

2 The language in brief

This chapter is enough pp to read the rest of the manual. pp is a Lisp-1: functions and variables share one namespace, and evaluation is inner-first and strict. The language reference appendix has the full detail of every form. Here is the working subset.

2.1 Values

pp has the atoms you expect: integers and floats, strings, the booleans `true` and `false`, `nil`, keywords like `:host`, and quoted symbols like `quote { name }`. Arithmetic and comparison are variadic. `if` is an expression: it returns a value.

hello.pp

```
# Every value has a content hash; arithmetic and comparison are variadic.
print(+ (1, 2, 3))
print(if 5 > 0 { "positive" } else { "negative" })
print(6 * 7)

$ pp hello.pp
```

```
6
"positive"
42
```

2.2 Bindings

`let` introduces local bindings. They are mutual: every binding can see every other, whatever order you write them in. For the ordinary top-down behaviour, where each binding sees only the ones above it, use `let*`.

lang-bindings.pp

```
# Bindings in a `let` are mutual: every binding sees every other, regardless
# of the order they are written.
print(let (y = x + 1, x = 1) { y })

# `let*` is the sequential form, where each binding sees only the ones above.
print(let* (x = 1, x = x + 1) { x })

$ pp lang-bindings.pp
```

```
2
2
```

Both forms bind names to the values of their right-hand sides. `def` does the same at the top level, and also defines functions.

2.3 Functions

A function definition and a function value are the same idea written two ways. `def` names one; `fn` produces one anonymously. There is no `loop` keyword: recursion is the loop. Both back ends eliminate tail calls, so recursion does not grow the stack.

lang-functions.pp

```
# `def` names a function; the last expression is its result.
def square(x) { x * x }
print(square(7))

# Functions are values. `fn` makes one without naming it.
print((fn(x) { x * x })(7))

# Recursion is the loop.
def fact(n) { if n = 0 { 1 } else { n * fact(n - 1) } }

print(fact(5))

$ pp lang-functions.pp
```

```
49
49
120
```

2.4 Types

Type annotations are optional. When present, `pp` checks them at the moment a function's body runs, not before. A mismatch names the offending value and location:

type-error.pp

```
# Type annotations are optional, and checked when the body runs. Passing a
# string where an int is annotated is caught – the manual shows the real error.
def inc(n: int): int { n + 1 }
inc("oops")

$ pp type-error.pp
```

```
pp: error: type mismatch: expected int, got "oops" at type-error.pp:3
```

2.5 Laziness

`delay` builds a thunk: a computation that has not run yet. `force` runs it and memoizes the result. The work happens at most once, however many times you force it.

lang-laziness.pp

```
# `delay` builds a thunk that has not run yet; `force` runs it, and the
# result is memoized – the work happens at most once.
let (t = delay(do {
# prints "working", then 20000
# prints 20000 – no re-run
  print("working")
  100 * 200
})) {
  print(force(t))
  print(force(t))
}
```

```
$ pp lang-laziness.pp
```

```
"working"
20000
20000
```

This is more than a convenience. Two identical thunks with the same inputs and environment are the same thunk: computed once, shared everywhere. Wrapping a computation in node (the next chapter) extends that sharing across separate runs of the program.

2.6 Effects and handlers

Side effects go through perform, which dispatches to the nearest enclosing handler. Handling an effect yourself needs no capability. The authority model in the chapter on capabilities governs only the effects that reach out and touch the world.

lang-effects.pp

```
# A side effect is a `perform` that dispatches to the nearest handler. This
# needs no capability: the authority model governs effects that touch the
# world, not ones you handle yourself.
print(with-handler(ask = fn(q) { 42 }) { perform ask("the answer?") })
```

```
$ pp lang-effects.pp
```

```
42
```

2.7 Config

Config is ambient data: dynamically scoped values that any code in the dynamic extent can read. It is deliberately not a capability. Config is information, capabilities are authority, and the manual keeps them apart.

lang-config.pp

```
# Config is ambient, dynamically-scoped data – distinct from capabilities,
# which are authority. A binding is visible to everything called within.
print(with-config({:host -> "db1"}) { config(:host) })
```

```
$ pp lang-config.pp
```

```
"db1"
```

That is the surface. Everything from here on is about what happens when you wrap a computation in node and give pp the authority to act on the world.

3 Nodes and content-addressing

The previous chapter ended on `delay` and `force`: a thunk runs at most once, and two identical thunks in one run are the same thunk. `node` extends that idea past the edge of a single run. You wrap a computation in `node`. Its result is content-addressed and written to a store on disk (`~/ .pp/store`). The next computation that asks for the same thing finds the answer already there, instead of recomputing it. That next computation might come later in this run, in a separate process, or on another machine sharing the store.

3.1 The node key

What makes two computations “the same thing” is a hash. A node’s key is

```
H(code || free-var value-hashes)
```

— the hash of the node’s code together with the values of the free variables that code references (this is LAW 20). Nothing else is in the key: not the ambient capabilities, not the config, not the handler stack, not the rest of the environment. Identical code fed identical inputs produces an identical key, and the same key is the same node.

Write the same node twice and `pp` treats them as one. The first `force` runs the body and stores the result. The second finds that result by key. A store hit does not replay whatever the body printed, so the side effect happens exactly once:

node-identity.pp

```
# A node is identified by its code and its inputs. Write the same node twice
# and it is the SAME node: the first force runs the body and stores the
# result; the second finds that result already in the store. A store hit does
# not replay the body's output (LAW 17), so "compiling" prints exactly once.
let a = node { print("compiling"); 6 * 7 }
let b = node { print("compiling"); 6 * 7 }
print(a)
print(b)

$ pp node-identity.pp
```

```
42
42
```

The key folds in the values of free variables, not the code that computed them. `x` and `y` below are built differently but are the same value. So the two nodes keyed on them are the same node, and the body runs once. A different input means a different key, and runs:

node-freevars.pp

```
# The node key folds in the VALUES of the free variables the body references,  
# not the code that produced them (LAW 20). `x` and `y` are computed  
# differently but are the same value, so `(squared x)` and `(squared y)` are  
# the same node – the body runs once. A different value is a different node.  
def squared(n) { node { print("squaring"); n * n } }  
let x = 5  
let y = 2 + 3  
print(squared(x))  
print(squared(y))  
print(squared(9))  
  
$ pp node-freevars.pp
```

```
25  
25  
81
```

This is why the key excludes the whole environment. Rebind an unrelated global and every node keyed on it is untouched. Change a value a node actually reads, and that node, and only that node, re-keys.

3.2 Caching across runs

Because the store is on disk, node results outlive the process that produced them. A delay memoizes within one run; a node memoizes across runs. Run a program, exit, then run it again. The second process serves the first's results without repeating the work:

node-reuse.sh

```
#!/bin/sh
# A node's result is content-addressed in ~/.pp/store, so it survives the
# process that produced it. A second, separate run of the same program finds
# the result already there and does not re-run the body. A throwaway store (a
# temp HOME) keeps this hermetic.
export HOME=$(mktemp -d)
cd "$HOME"

cat > build.pp <<'PP'
def compile() { node { print("compiling greeter.o"); 6 * 7 } }
print(compile())
PP

echo '$ pp build.pp      # first process: the node body runs'
"$PP" build.pp

echo
echo '$ pp build.pp      # second process: served from the store, body skipped'
"$PP" build.pp

rm -rf "$HOME"
```

```
$ pp build.pp      # first process: the node body runs
"compiling greeter.o"
42

$ pp build.pp      # second process: served from the store, body skipped
42
```

The body prints on the first run and is silent on the second. The second process never entered the body. This is the same collapse force gives you in memory, made durable: equal computations run once, no matter how many processes ask.

3.3 Pure compute needs no capability

None of these examples were granted anything. `node-identity` and `node-freevars` run with a bare `pp` file, `pp`, no `--grant`, because arithmetic touches nothing outside itself. Capabilities govern authority over the world, and the chapter on capabilities covers them. A computation that only computes needs none, and caching it needs none either. The moment a node reads the world — a file, the environment, a subprocess — the cache has to account for what it saw. That accounting is the subject of the next chapter.

4 The store and verifying traces

A node's key decides identity: which computation this is. It does not decide whether a cached result is still good. That is a separate question, and `pp` answers it in a way that never trusts a clock.

4.1 The store

Everything a node produces lives under `~/.pp/store`, in three content-addressed parts:

store-layout.sh

```
#!/bin/sh
# The store has three content-addressed parts. A temp HOME keeps it throwaway.
export HOME=$(mktemp -d)
cd "$HOME"

cat > build.pp <<'PP'
let artifact = node { blob(slurp("input.txt")) }
print(artifact)
PP
echo hello > input.txt

echo '$ pp build.pp --grant fs:$HOME:ro'
"$PP" build.pp --grant "fs:$HOME:ro"

echo
echo '$ ls ~/.pp/store'
ls "$HOME/.pp/store"

rm -rf "$HOME"
```

```
$ pp build.pp --grant fs:$HOME:ro
"blob:5891b5b522d5df086d0ff0b110fbd9d21bb4fc7163af34d08286a2e846f6be03"

$ ls ~/.pp/store
blobs
locks
objects
traces
VERSION
```

- objects/ — result values, each file named by the hash of the value it holds.
- traces/ — one file per node key, holding the set of traces recorded for that node (see below).
- blobs/ — raw bytes ingested by `blob(...)`, named by their hash. The program above turned a file's contents into a blob: reference, and the bytes landed here.

(VERSION stamps the store format and locks/ guards concurrent writers. Neither is content you address directly.)

4.2 A hit is a verified trace, not a timestamp

When a node runs, pp records every observation it makes of the world as a (cell, hash) pair: each file it reads, each config value, each subprocess tool. That record is a verifying trace: not “this ran at 3pm” but “this is what the node saw.”

On the next force, pp does not check a modification time. It re-observes every cell in the trace, and serves the stored result only if every observed hash still matches. A changed input fails the match and forces a recompute. An unchanged input matches and hits. One key holds a set of traces. So a node that ran under two different toolchains or platforms keeps a valid trace for each, and reverting a file re-matches its older trace. A cache hit is therefore always re-checked against the current world. It can never serve stale data, because “stale” means “some observed cell no longer matches”, and that is exactly what the hit test checks.

pp checks authority at the same moment. It serves a stored result only if the caller holds capabilities covering every file cell in the trace's read closure (LAW 23b). Reads propagate to enclosing nodes, so

the check is transitive: a narrow caller cannot launder a broad read through a cached aggregator. A capability denial is never cached. Authority gates access to a result, it does not rename it.

4.3 pp why

`pp why file.pp` runs the program and reports, to `stderr`, the fate of every node force: a first-build miss, a verified hit, or a stale trace naming the cell that changed. Here a build reads a file, the same run repeated hits, then you edit the input and the node re-runs:

store-rebuild.sh

```
#!/bin/sh
# A cache hit is decided by re-checking what the node observed, not by a
# timestamp. `pp why` reports each node's fate. A temp HOME keeps the store
# throwaway; ROOT is the canonicalized HOME, stripped from output so the
# transcript is machine-independent (the cell ids are absolute paths).
export HOME=$(mktemp -d)
ROOT=$(cd "$HOME" && pwd -P)
cd "$HOME"

cat > build.pp <<'PP'
let greeting = node { slurp("greeting.txt") }
print(greeting)
PP
echo hello > greeting.txt

echo '$ pp why build.pp --grant fs:$HOME:ro    # first build: a miss'
"$PP" why build.pp --grant "fs:$HOME:ro" 2>&1 | sed "s|$ROOT/||g"

echo
echo '$ pp why build.pp --grant fs:$HOME:ro    # input unchanged: a hit'
"$PP" why build.pp --grant "fs:$HOME:ro" 2>&1 | sed "s|$ROOT/||g"

echo
echo goodbye > greeting.txt
echo '$ pp why build.pp --grant fs:$HOME:ro    # input edited: re-run'
"$PP" why build.pp --grant "fs:$HOME:ro" 2>&1 | sed "s|$ROOT/||g"

rm -rf "$HOME"
```

```
$ pp why build.pp --grant fs:$HOME:ro    # first build: a miss
[why] node f1b010a0e48c: miss – no stored trace (first build)
"hello\n"

$ pp why build.pp --grant fs:$HOME:ro    # input unchanged: a hit
[why] node f1b010a0e48c: hit – ok trace verified (1 cells)
"hello\n"

$ pp why build.pp --grant fs:$HOME:ro    # input edited: re-run
[why] node f1b010a0e48c: trace 1/1 stale – file:greeting.txt changed
[why] node f1b010a0e48c: miss – no stored trace usable
"goodbye\n"
```

The node key (f1b010a0e48c here) is stable: it is a hash of the code, which did not change. What changed is the world the trace was verified against. The first run has nothing to verify against and misses. The second re-observes the one file cell, finds it unchanged, and hits. Editing the file makes that cell's hash no longer match. The stored trace goes stale, no other trace in the set is usable, and the node recomputes.

4.4 Auditing the cache

Two flags let you check the cache instead of trusting it.

--no-cache skips cache reads, so every node recomputes, while still writing fresh results and traces. Use it to force a clean build. A node that would otherwise hit reports its miss as `cache reads disabled` (--no-cache).

--check re-runs each missed node's body and compares the new result hash against the stored one. Matching hashes confirm the node is deterministic. A divergence flags the node as volatile and fails the run. This is how you catch a computation that claims to be a pure function of its inputs but is not, the one thing content-addressed caching cannot tolerate.

5 Building with pp

A build system is a graph of tool runs with caching: compile each source to an object, link the objects, don't redo work whose inputs haven't changed. pp already has that graph. A node is a cached computation keyed by its code and inputs; the store keeps its result. Point a node at a compiler and you have a build system — no new machinery, just the effect that runs a process and the primitives that move bytes into the store.

This chapter builds a real C program that way. It then shows the three things a build system has to get right: a null rebuild does nothing, a one-line edit recompiles only what depends on it, and deleting the outputs restores them without re-running a single tool.

5.1 A compilation unit is a node

Running a tool is a side effect, so it goes through `perform`, and it needs authority — the process capability from the chapter on capabilities. The plain form is `run`; the one you build with is `run-dep`:

```
perform run-dep("greet.d", "cc", "-MD", "-MF", "greet.d", "-O0", "-c",
  "/abs/src/greet.c", "-o", "greet.o")
```

`run-dep` runs the command exactly like `run` — the first extra argument names a Makefile-style depfile the tool is told to write (`cc -MD -MF greet.d`). The command runs in the node's own sandbox directory, so relative outputs like `greet.o` land in scratch that no other node can see. That output is not yet a value; you pull it into the store with `blob` and `slurp`:

```
blob(slurp("greet.o"))
```

`slurp` reads the sandbox file; `blob` writes its bytes into the content-addressed store and returns a short `"blob:<hash>"` reference. Wrap the whole thing in a node and a compilation unit becomes a cached value: the same source and the same command hash to the same node, so the object is computed once and shared by every later run.

5.2 Depfiles refine the trace

A node records which cells it read, and re-runs only when one of them changes. The question is what a `cc` invocation reads. The conservative answer — the whole source tree — would recompile everything on any edit. That is what `run-dep` avoids. After the command exits, pp parses the depfile the tool wrote. It refines the node's trace to exactly the files the compiler actually opened. `greet.c` includes `greet.h`, so `greet.o`'s node depends on those two files and nothing else. `main.c` includes the same header, so it depends on `main.c` and `greet.h`.

That refinement is the whole basis of incrementality below: editing `greet.c` touches `greet.o`'s trace but not `main.o`'s, so only `greet.o` is rebuilt. A tool that writes no depfile falls back to the coarse whole-tree dependency — sound, just less precise. The trust is explicit and per-call: you chose `run-dep`.

5.3 Blobs and materialization

A node returns a blob: reference, not bytes. The program's final value is a map from output path to reference — the build tree you want on disk:

```
{"greet.o" -> "blob:...",  
  "main.o" -> "blob:...",  
  "prog"   -> "blob:...:x"}
```

The trailing :x marks an executable. You hand that map to the reconciler with `--reconcile <root>`. It diffs the map against what is already under <root> and materializes the difference from the store, writing a file only when its hash differs from the one on disk. The reconciler and its diff-and-converge loop are the subject of a later chapter; here it is just the step that turns the build tree into files. Because outputs are addressed by content, materializing them never re-runs a tool: the bytes are already in the store.

5.4 The incremental story

Here is the whole build as one shell transcript. The program compiles two translation units and links them; the grants give it read access to the sources, write access to the build directory, and process authority. Rather than show timings, which aren't reproducible, it counts external processes straight from the store's journal — the same journal that makes “zero processes” a fact you can check rather than a claim.

build-incremental.sh

```
#!/bin/sh
# Building a tiny C program with pp, then the incremental story. A throwaway
# store (temp HOME) keeps this hermetic; output is counts, never timings, so
# it is byte-reproducible.
export HOME=$(mktemp -d)
cd "$HOME"
mkdir -p src build

# --- a two-file C program: main() calls greet() ---
cat > src/greet.h <<'EOF'
void greet(void);
EOF
cat > src/greet.c <<'EOF'
#include "greet.h"
#include <stdio.h>
void greet(void) { printf("hello from pp\n"); }
EOF
cat > src/main.c <<'EOF'
#include "greet.h"
int main(void) { greet(); return 0; }
EOF

# --- the build: one node per translation unit, then a link node ---
# compile: run-dep records the EXACT files cc read (the .d depfile), then blob
# ingests the .o into the store. link: write each object into the node's
# sandbox, link them, blob the result. The program's value is the desired
# build tree: {relpath -> blob-ref}, ":x" marking the executable.
cat > build.pp <<EOF
def each(f, lst) { if nil?(lst) { nil } else { f(car(lst)); each(f, cdr(lst)) } }
def foldl2(f, acc, lst) { if nil?(lst) { acc } else { foldl2(f, f(acc, car(lst)), cdr(lst)) } }
def zip2(a, b) {
  if nil?(a) { nil } else { cons(cons(car(a), car(b)), zip2(cdr(a), cdr(b))) }
}
def compile(name) {
  node {
    perform run-dep(string-append(name, ".d"), "cc", "-MD", "-MF", string-
append(name, ".d"), "-O0", "-c", string-append("$HOME/src/", name, ".c"), "-o",
string-append(name, ".o"))
    blob(slurp(string-append(name, ".o")))
  }
}

def link(objs) {
  force(node { each(
fn(o) { perform write-file(string-append(car(o), ".o"), blob-get(cdr(o))) }, objs)
do {
  perform write-file("link.d", "prog: ")
  do {
    perform run-dep("link.d", "sh", "-c", "cc -o prog greet.o main.o")
    blob(slurp("prog")) } } }) }
let (names = list("greet", "main"), objs = zip2(names, force-deep(map(compile,
names))), prog = link(objs)) {
  foldl2(
```

```
fn(o) { perform write-file(string-append(car(o), ".o"), blob-get(cdr(o))) },
"prog", string-append("prog", blob-ref(build:prog)) }
--reconcile build build.pp
}
```

```
EOF
-- cold build
processes: 3
```

```
C file: /tmp/268d4205-2025-411e-b011-000000000000/build:prog
C file: /tmp/268d4205-2025-411e-b011-000000000000/build:prog
```

Read the four sections in order. The cold build runs three processes — two compiles and one link — and materializes three files. The second build changes nothing: every node is a store hit, so zero processes run and the reconciler writes nothing (`create=0 update=0 delete=0`). This is the null rebuild, and it does no work because there is no work whose inputs changed.

Editing `src/greet.c` changes one file. Its depfile refined `greet.o`'s trace to `greet.c` and `greet.h`, so `greet.o` rebuilds; `main.o`'s trace didn't include `greet.c`, so it is served from the store untouched. The transcript confirms it: `recompiled: greet.c`, two new processes (the one compile plus the relink). This is the property that makes an incremental build worth having — work is proportional to what changed, not to the size of the project.

Finally, `rm -rf build` deletes every output. The rebuild runs zero processes: the reconciler finds an empty tree, diffs it against the same build map, and materializes all three files straight from the store. The restored binary is byte-identical to the one deleted. Outputs are a cache of the store, and the store is the source of truth — so losing them costs a copy, never a compile.

6 Capabilities and secrets

A pure computation needs no permission: `pp` will add, compare, and build lists all day without asking. The moment a computation reaches out and touches the world — reads a file, runs a process, makes a request — it needs a capability: an unforgeable token that says it is allowed to. No capability, no effect. This is not a sandboxing layer bolted on the side; it is the only way a world-touching effect ever runs.

Here is a program that tries to read a file with no authority at all:

cap-read.pp

```
# Reading a file reaches out and touches the world, so it needs a filesystem
# capability. This program is granted nothing, so the read is refused before
# it happens — the check is on authority, not on whether the file exists.
print(perform read-file("/etc/hostname"))
```

```
$ pp cap-read.pp
```

```
pp: error: read-file: capability error: no read access for /etc/hostname at cap-read.pp:4
```

The read is refused before it happens. The check is on the authority, not on whether the file exists or what it contains — `pp` never looked. The error names the effect, the missing access, and where the attempt was made.

6.1 Capabilities enter only at the root

There is no expression that creates authority. You cannot write `filesystem("/", :rw)` — `filesystem` is an unbound symbol, not a constructor. The sole mint is the command line: `--grant` hands the program a capability as it starts, and nothing inside the language can manufacture one. Grant the read above a covering `filesystem` capability and the same effect succeeds:

cap-grant.sh

```
#!/bin/sh
# The same read-file effect that fails without authority succeeds once the
# root grants a covering capability. Capabilities enter ONLY here, on the
# command line: there is no expression that mints one. A throwaway HOME with
# a file under it keeps the run hermetic; we cd into it and use relative
# paths so nothing machine-specific reaches the output.
export HOME=$(mktemp -d)
mkdir -p "$HOME/etc"
printf 'dbl.internal' > "$HOME/etc/hostname"
cd "$HOME"

cat > read-hostname.pp <<'PP'
print(perform read-file("etc/hostname"))
PP

echo '$ pp read-hostname.pp                # no grant: refused'
"$PP" read-hostname.pp 2>&1

echo
echo '$ pp --grant fs:etc:ro read-hostname.pp    # covering grant: allowed'
"$PP" --grant fs:etc:ro read-hostname.pp 2>&1

rm -rf "$HOME"
```

```
$ pp read-hostname.pp                # no grant: refused
pp: error: read-file: capability error: no read access for etc/hostname at read-hostname.pp:1

$ pp --grant fs:etc:ro read-hostname.pp    # covering grant: allowed
"dbl.internal"
```

A grant is a scheme, a target, and — for the filesystem — a mode:

- `--grant fs:<path>:ro` (or `:rw`, `:wo`) — read, write, or read-write access to a directory and everything under it. Scope matching is by path component on the canonicalized full path, so `fs:/tmp:ro` covers `/tmp/x` but never `/tmpevil`.
- `--grant net:<host>` or `net:<host>:<port>` — network access to a host, and optionally one port; `net:*` wildcards the host.
- `--grant secret:<path>` — a sealed read of the files under a path (below).
- `--grant process` — the authority to run subprocesses.

Pass `--grant` more than once to hold several capabilities at once.

6.2 User code narrows, never widens

Code holds capabilities, passes them around, and attenuates them — but only ever downward. `current-capabilities()` reifies the ambient authority as of the call; it is an observation of the ceiling every `perform` already checks against, not a mint. Two forms narrow it:

- `cap-restrict(cap, scope)` — restricts `cap` to a sub-scope, optionally to a narrower mode, for example `cap-restrict(cap, "src", :ro)`. Asking for a mode wider than `cap` already grants at that scope is a `Capability_error`, never a silent widen.
- `cap-compose(a, b, ...)` — unions capabilities the code already holds; it can only ever produce authority already present in its arguments.

`with-caps(cap-expr) { body }` then replaces the ambient authority with `cap-expr` for the extent of `body`. A subset check against the current ambient gates it. So a narrowing composes even when some other binding still lexically holds a broader value. The narrowing is restored when `body` returns, tail-calls, or raises. There is deliberately no union-with-ambient form: the instant capability values exist, unioning with the ambient would be a widening backdoor.

Authority is a ceiling, re-checked at every `perform`. It is never part of a value's identity: widening the grant does not rename or invalidate anything, and narrowing it does not.

6.3 Authority gates the cache, transitively

A cache hit is a message from a past run to this one, so authority has to gate it too — not just live effects. When `pp` considers serving a node's stored result, it requires the caller's capabilities to cover the transitive read closure of that result: every cell the node read, and recursively every cell its child nodes read. A caller scoped to `src/` is not served a hit on a node whose closure touched `/etc/passwd`. That holds even though the expensive work is already done and sitting in the store. Being handed the result would tell it that the read happened and what it produced. A capability denial is not memoized, so granting the authority later still yields the hit.

This is also why a node may not carry a capability across its own boundary in either direction: a node whose free variable is (or contains) a capability is rejected before it is keyed, and a node whose result is (or contains) one is rejected before it is stored. Authority that rode a cached result out to a narrower caller would defeat the whole check.

6.4 Secrets: sealed cells

A secret is authority-adjacent but a different problem: you want a computation to use a value without that value's bytes leaking into the terminal or, worse, into the content-addressed store, where any later run could read them back. A `secret: grant` handles this. A read covered by `secret: —` and not also by a plain `fs: grant` over the same path — returns a sealed value instead of an ordinary string:

cap-secret.sh

```
#!/bin/sh
# A secret is granted with `secret:` instead of `fs:`. A read covered by a
# secret grant (and NOT by a plain fs grant) yields a SEALED value: it prints
# redacted, its bytes are pinned in memory only and never written to the
# content-addressed store, yet a hash of those bytes still rides in the node's
# trace – so rotating the secret invalidates exactly its observers. Hermetic:
# throwaway HOME, a secret file under it, relative paths.
export HOME=$(mktemp -d)
mkdir -p "$HOME/vault"
printf 'hunter2-swordfish' > "$HOME/vault/token"
cd "$HOME"

# (1) A sealed value prints redacted – the bytes never reach the terminal.
cat > show.pp <<'PP'
print(slurp("vault/token"))
PP
echo '$ pp --grant secret:vault show.pp'
"$PP" --grant secret:vault show.pp 2>&1

# (2) Derive from the secret inside a node (its length), which populates the
# store, then scan the whole store for the plaintext.
echo
cat > derive.pp <<'PP'
print(force(node { string-length(unseal(slurp("vault/token"))) }))
PP
echo '$ pp --grant secret:vault derive.pp      # length of the secret'
"$PP" --grant secret:vault derive.pp 2>&1

echo
echo '$ grep -r hunter2-swordfish ~/.pp/store # the bytes are not there'
if grep -rq 'hunter2-swordfish' "$HOME/.pp/store" 2>/dev/null; then
  echo 'FOUND (leak)'
else
  echo '(no match – secret bytes absent from the store)'
fi

rm -rf "$HOME"
```

```
$ pp --grant secret:vault show.pp
#<sealed>

$ pp --grant secret:vault derive.pp      # length of the secret
17

$ grep -r hunter2-swordfish ~/.pp/store # the bytes are not there
(no match – secret bytes absent from the store)
```

Three things hold at once. The sealed value prints redacted as `#<sealed>`, so it cannot leak by being printed. Its bytes are pinned in memory only and are never written to the store, so a scan of the whole store never finds them — the `grep` above comes up empty even though a node ran and populated the store. And a hash of those bytes does ride in the node's trace, so the cache stays correct: rotate the secret and exactly the nodes that observed it are invalidated, while their siblings keep hitting.

Deriving from a secret inside a node — the length, above — produces ordinary data, because the derived value is no longer the secret. The one explicit way to turn a sealed value back into a plain string is

`unseal(v)`; there is no implicit dataflow tainting. Unsealing is a deliberate, greppable boundary, so the confidentiality of a secret ends exactly where the program says it does. If a path is covered by both a `secret:` and an `fs: grant`, the ordinary filesystem behaviour wins: the deployment that granted plain access is saying “not secret here”.

7 Modules and islands

A module is a block of code whose definitions become a value. That is the whole idea. Packaging is not a separate mechanism layered over the language. It is what happens when you take the names a block defined and hand them back as something you can bind, pass, and import. Islands extend the same value across machines by pinning it to a content hash.

7.1 Modules

`module { ... }` evaluates its body in a fresh scope and produces a value carrying the names it defined. `import(m)` forces that value and merges those names into the current scope:

mod-module.pp

```
# (module ...) evaluates its body in a fresh scope and produces a value whose
# fields are the names it defined. (import ...) forces that value and merges
# those names into the current scope. Here the module is bound to `geo`, then
# imported, so `area` and `pi` become callable.
let (geo = module {
  let pi = 3.14159
  def area(r) { pi * (r * r) }
}) {
  import(geo)
  print(area(10))
}

$ pp mod-module.pp
```

```
314.159
```

The module’s body sees only itself, not the scope around it, so a module is a sealed unit: what escapes is exactly what it defined, and only once you import it. You can hold a module value, pass it to a function, or import it conditionally — it is an ordinary value until the moment you pull its names in.

7.2 Loading from files

Two forms bring a file’s definitions into a program, and they differ in whether the file gets its own scope:

- `load("f.pp")` evaluates the file’s forms in the current scope — its definitions merge in directly, as if you had typed them here.
- `load-module("f.pp")` evaluates the file isolated, in its own fresh scope, and returns its exports as a module value — which you then import. The file’s internal scope never leaks into yours.

mod-load-module.sh

```
#!/bin/sh
# Two ways to bring a file's definitions in. `load` MERGES a file into the
# current scope (its defs become directly visible). `load-module` loads the
# same file ISOLATED and returns its exports as a value, which you then
# `import` — so the file's own scope never leaks in. Hermetic throwaway HOME.
export HOME=$(mktemp -d)
cd "$HOME"

cat > mathlib.pp <<'PP'
def square(x) { x * x }
def cube(x) { x * square(x) }
PP

cat > use-load.pp <<'PP'
# load: mathlib's defs merge into this scope.
load("mathlib.pp")
print(square(9))
PP
echo '$ pp use-load.pp          # load merges definitions in'
"$PP" use-load.pp 2>&1

echo
cat > use-load-module.pp <<'PP'
# load-module: mathlib loads isolated; its exports come back as a value.
let (m = load-module("mathlib.pp")) {
  import(m)
  print(cube(3))
}
PP
echo '$ pp use-load-module.pp  # load-module returns exports to import'
"$PP" use-load-module.pp 2>&1

rm -rf "$HOME"
```

```
$ pp use-load.pp          # load merges definitions in
81

$ pp use-load-module.pp  # load-module returns exports to import
27
```

Reach for `load` when a file is a fragment of the current program, and `load-module` when it is a self-contained unit you want to keep at arm's length. Either way, the read is the loader's own — it runs under the interpreter's runtime authority, bounded to your source roots, and is exempt from the capability accounting of the previous chapter. Editing a loaded file still invalidates the nodes that loaded it: the load is recorded in their traces even though it costs no user capability.

7.3 Islands

An island is a module that lives elsewhere — another directory, a git repository, a URL — referenced by URI and pinned inline by the content hash of its source tree:

```
island("file:./lib", "e5d7b0f8...")
island("github:owner/repo#main", "a1b2c3...")
```

The pin is part of the code. Because it folds into the expression's hash, island identity is structural: there is no lockfile and no hidden resolution step — a pinned island form denotes the same bytes wherever you paste it. The ref after # lives in the URI and matters only when fetching; the pin argument is always the 64-hex content hash of the tree.

Resolution never touches the network. The pin names an immutable tree in the island cache under `~/.pp/islands/src/<pin>/`, whose `entry.pp` is the module root. `pp` verifies the cached tree against the pin on every resolve. A mismatch is a hard error, not a silent reuse. An unpinned island form is a hard error too, naming the fix. Deriving the pin, and fetching a `git:` or `github:` tree the first time, is the one impure step, and it is opt-in. `pp --update` re-resolves each island and rewrites the pins in your source, while `git:/github:` fetching happens only under `--fetch-islands`. With neither enabled, evaluation stays hermetic.

The workflow, end to end, on a local file: island:

mod-island.sh

```
#!/bin/sh
# An island is a module referenced by URI and pinned INLINE by the content
# hash of its source tree. Resolution never touches the network: the pin names
# an immutable tree in the island cache. An unpinned form is a hard error that
# names the fix; `pp --update` derives the pin and writes it into the source;
# the pinned form then evaluates the tree's entry.pp as a module. Hermetic
# throwaway HOME (which also holds the island cache).
export HOME=$(mktemp -d)
mkdir -p "$HOME/lib"
printf 'def greet(who) { string-append("hello, ", who) }\n' > "$HOME/lib/entry.pp"
cd "$HOME"

cat > app.pp <<'PP'
let (m = island("file:./lib")) {
  import(m)
  print(greet("world"))
}
PP

echo '$ pp island-pins app.pp    # the form is unpinned'
"$PP" island-pins app.pp 2>&1

echo
echo '$ pp --update app.pp      # derive the pin and write it into the source'
"$PP" --update app.pp 1>/dev/null

echo
echo '$ cat app.pp              # the pin is now part of the code'
cat app.pp

echo
echo '$ pp app.pp               # resolve the pinned tree and run it'
"$PP" app.pp 2>&1

rm -rf "$HOME"
```

```
$ pp island-pins app.pp    # the form is unpinned
file:./lib (unpinned)

$ pp --update app.pp      # derive the pin and write it into the source

$ cat app.pp              # the pin is now part of the code
let (m = island("file:./lib", "da7539569962c864b52f38eacb47fa03b6d4304f7d09ac5bed134c15a54f8dbf")) {
  import(m)
  print(greet("world"))
}

$ pp app.pp               # resolve the pinned tree and run it
"hello, world"
[update] app.pp: 1 pin(s) updated, 0 skipped
```

`pp island-pins` reports the forms and whether each is unpinned, cached, or tampered; `pp --update` fills in a missing pin rather than making you compute the hash by hand, and refuses to touch a form it cannot rewrite unambiguously rather than half-writing your file. Once the pin is in place, the island is just another module: `import` its value and call what it exports.

8 Domains and the reconciler

A node computes a value. But a build has to change the world: write files, start processes. pp will not let you do that by mutating it directly. Instead you compute a desired state: a pure, hashable value that says what the world should contain — $\{\text{path} \rightarrow \text{content}\}$, $\{\text{service} \rightarrow \text{spec}\}$. A domain observes the world, diffs it against your desired value, and applies the minimal change to converge them. This is React’s contract verbatim. You never touch the DOM. You return the desired DOM and the reconciler applies the diff. Here the world is the filesystem and the process table.

The payoff is that convergence follows from the value, not from a script you maintain. Your desired state is a pure function of input cells. So pp caches it, re-derives it cheaply when an input changes, and re-observes reality rather than trusting a state file. Drift is just a difference the next pass erases — a file someone edited by hand, a service that died.

8.1 A domain is observe / diff / apply

A domain is three ordinary pp functions registered over a namespace of cells:

- `observe : () → value` reads the current world. It runs fresh every pass, never cached.
- `diff : (observed, desired) → plan` computes what to change. It is pure: it runs under an empty capability set and cannot touch the world. Its result is cached, keyed on the diff code plus the observed and desired values.
- `apply : plan → nil` performs the plan’s changes, under the domain’s own write capability.

Core wraps every domain’s apply in one discipline: a journalled intent/done bracket, atomic writes, and verify-after-write. Core re-observes and re-diffs once the apply returns. A plan that still has work left is a hard error. A domain is the single writer for its namespace. So there are no write-write races, and your code needs no ordering discipline. One rule enforces this: stratification. The desired state may not read the domain’s own cells, or reconcile would loop forever. Reading your output tree to decide your output tree is refused.

The trusted mechanics that actually touch the world are a small set of core primitives: atomic temp-file-plus-rename, fork/exec/reap, a per-domain key-value store (tree-observe, materialize-file, remove-file, proc-spawn, and the like). Everything else is a pp library. The filesystem and process domains ship as `stdlib/domain-fs.pp` and `stdlib/domain-proc.pp`. This is real pp source you can read. It holds all the policy over the core-enforced protocol: what counts as a create versus an update, when to restart a service. There is no privileged reconciler engine hidden in the runtime. A domain is a library.

8.2 Reconciling a filesystem

`pp --reconcile R00T prog.pp` auto-loads the fs domain and registers it with a write capability narrowed to R00T. It takes the program’s final value as the desired state of the tree under R00T: a $\{\text{relative-path} \rightarrow \text{content}\}$ map. It diffs that against the real directory by content hash and materializes missing and changed files atomically. It also deletes files under R00T that the map does not mention. The domain is the single writer, and the write grant is your consent to that authority. An fs write grant over R00T is required. Without it, nothing is written.

The transcript below reconciles a two-file desired state into a fresh root, so both files appear. It then introduces drift by hand, deleting one file and editing another, and reconciles again. The second pass reports exactly one create and one update, and the tree is restored. The summary line pp prints names the counts per kind. Its absolute root path is filtered to R00T so the output is machine-independent.

domain-reconcile.sh

```
#!/bin/sh
# A domain converges the world to a DESIRED state you return as a value. The
# fs domain (stdlib/domain-fs.pp) takes a {relative-path -> content} map and
# makes a directory tree match it: creating, updating, and deleting files.
# Hermetic: a throwaway HOME (so the store and journal are fresh) and a root
# under it. The reconcile summary's absolute root path is filtered to ROOT.
export HOME=$(mktemp -d)
ROOT="$HOME/site"

cat > "$HOME/site.pp" <<'PP'
{"index.html" -> "<h1>pp</h1>\n", "conf/app.txt" -> "mode=prod\n"}
PP

filter() { sed "s#root=[^ ]*#root=ROOT#"; }

echo '$ pp --grant fs:ROOT:rw --reconcile ROOT site.pp # first pass: create'
"$PP" --grant "fs:$ROOT:rw" --reconcile "$ROOT" "$HOME/site.pp" 2>&1 | filter
( cd "$ROOT" && find . -type f | sort )
echo "conf/app.txt: $(cat "$ROOT/conf/app.txt")"

echo
echo '# Introduce drift: delete one file, corrupt another.'
rm "$ROOT/index.html"
echo 'mode=DEV' > "$ROOT/conf/app.txt"

echo '$ pp --grant fs:ROOT:rw --reconcile ROOT site.pp # second pass: restore'
"$PP" --grant "fs:$ROOT:rw" --reconcile "$ROOT" "$HOME/site.pp" 2>&1 | filter
( cd "$ROOT" && find . -type f | sort )
echo "conf/app.txt: $(cat "$ROOT/conf/app.txt")"

rm -rf "$HOME"
```

```
$ pp --grant fs:ROOT:rw --reconcile ROOT site.pp # first pass: create
[reconcile:fs] root=ROOT create=2 update=0 delete=0
./conf/app.txt
./index.html
conf/app.txt: mode=prod

# Introduce drift: delete one file, corrupt another.
$ pp --grant fs:ROOT:rw --reconcile ROOT site.pp # second pass: restore
[reconcile:fs] root=ROOT create=1 update=1 delete=0
./conf/app.txt
./index.html
conf/app.txt: mode=prod
```

Nothing in `site.pp` describes how to converge. There is no “if missing, create” logic. It states the desired contents, and the domain works out the difference. Reverting `conf/app.txt` and restoring the deleted `index.html` are the same mechanism, driven entirely by the diff.

Desired contents may be inline strings, as here, or `blob:<hash>` references into the content-addressed store — a compiled artifact ingested with `blob`. A blob reference diffs by hash without loading its bytes. So `rm -rf` on the tree restores from the store with zero tool re-runs when the desired-state nodes hit.

8.3 Watching, and other domains

`pp --watch --reconcile ROOT prog.pp` runs the program, reconciles, then polls the observed cells and re-runs on any change. This is a controller loop. Every registered domain is re-observed and re-

applied on each tick, whichever cell changed. This is cheap when nothing moved, because the plan cache turns a no-op pass into a hit. An externally deleted file or a drifted config is caught within one poll interval.

The process domain is the same protocol over a different world. `pp --supervise prog.pp` (usually with `--watch`) takes a `{service-name → spec}` map and keeps the process table matching it. It starts missing services, stops removed ones, and restarts a service whose spec changed. It also restarts one killed out from under it, within a poll interval. It requires `--grant process` and refuses stratification on its own `proc: cells`, exactly as the `fs` domain does. A from-scratch third-party domain is anything with an `observe/diff/apply` triple registered via `register-domain`. It gets the same journal bracket, plan cache, `verify-after-write`, and stratification for free. The protocol is generic, not filesystem-shaped.

8.4 Fenced effects

Convergence covers only idempotent change: applying a desired state twice is the same as applying it once. Some actions are not like that, such as sending an email or charging a card. Those may not appear in a node at all. A node is cache-replayable and must never carry an irreversible action. The scripting form `fenced(KIND, SPEC)` registers such an action for the reconciler to run after all convergent work, at most once per pass. It carries an intent/done journal. So a crash mid-action is recovered by policy (`--fenced-policy retry | abort | ask`) rather than by a silent retry that might double-charge. The desired-state law tames the convergent world. Fenced effects are the named carve-out for the part that cannot be converged.

9 Distribution

Here is the thesis this whole manual has been building toward: there is no local/remote distinction at the language surface. There is no `remote-eval`, no placement annotation, no node-pinning form, and there never will be. `force` is the only execution primitive. Where a force runs is a scheduler decision, chosen with a flag, never written into the program. A program is byte-identical whether it runs on one core, eight, or a cluster. Microservices exist because no mainstream language lets you say “evaluate this elsewhere and flow the result back”. The moment location becomes syntax, every caller hard-codes topology, and you have rebuilt the deployment boundary inside the language.

`--schedule serial | parallel:N | race:N | remote:<member>` selects the policy. `parallel:N` forks workers across cores. `race:N` runs `N` redundant copies and takes the first. `remote:<member>` ships the work to another machine. These are the same feature at different fan-out. “Run on `N` machines and take the first” is a handler swap, not an infrastructure project. The scheduler is result-transparent: it changes only where and when work runs, never the answer. So it appears in no cache key and no trace. `pp --check` proves this. It re-runs any non-serial policy forced serial against the same store and fails on any hash mismatch.

9.1 Identity is location-independent

Location can be a scheduler’s business, not the program’s, because identity is a content hash. A node’s key is a hash of its code and its inputs’ values, never of where or when it ran. Two machines that evaluate the same node therefore compute the same key and a byte-identical result, with no coordination. A result is then moved by that hash. It is published into a shared store and pulled by name, then re-hashed on receipt so a corrupted or tampered artifact is refused. The content address is the integrity check.

The transcript below uses two throwaway HOMEs as two machines, each with its own `~/ .pp` store, and a shared directory as the transport. Machine A builds a node. Machine B, given the same program,

computes the identical node key without ever talking to A. Then A publishes the result object by hash and B pulls it, byte-identical. Flipping one byte in transit makes B reject it.

dist-by-hash.sh

```
#!/bin/sh
# Distribution in pp is not a language feature: identity is a content hash, so
# a computation has the SAME name on every machine, and a result is moved by
# that hash. Two throwaway HOMEs stand in for two machines (each has its own
# ~/.pp store); a shared directory is the transport (a plain local-dir
# loopback – the ssh transport is stubbed, STATUS.md). Content hashes are
# deterministic, so they are reproducible output.
export HOME=$(mktemp -d)
A="$HOME/machine-a"; B="$HOME/machine-b"; SHARED="$HOME/shared"
mkdir -p "$A" "$B" "$SHARED"

cat > "$HOME/node.pp" <<'PP'
print(force(node {
  perform log("building greeter.o")
  6 * 7
}))
PP

echo '$ pp node.pp' # on machine A: the node body runs'
HOME="$A" "$PP" "$HOME/node.pp" 2>&1
KEY=$(ls "$A/.pp/store/traces")
RH=$(grep -oE '"[0-9a-f]{64,}"' "$A/.pp/store/traces/$KEY" | head -1 | tr -d '"')
echo "node key: $KEY"
echo "result hash: $RH"

echo
echo '$ pp node.pp' # on machine B, independently: SAME node key'
HOME="$B" "$PP" "$HOME/node.pp" 2>&1
KEY_B=$(ls "$B/.pp/store/traces")
[ "$KEY_B" = "$KEY" ] && echo "machine B computed the same key: yes"

echo
echo '# A publishes the result object by hash; B pulls it by that hash.'
echo '$ pp --transport-push object <hash> shared' # on A'
HOME="$A" "$PP" --transport-push object "$RH" "$SHARED" >/dev/null 2>&1
echo '$ pp --transport-pull object <hash> shared' # on B, re-hash-verified'
HOME="$B" "$PP" --transport-pull object "$RH" "$SHARED" 2>&1
diff -q "$A/.pp/store/objects/$RH" "$B/.pp/store/objects/$RH" >/dev/null \
&& echo "byte-identical across the transport: yes"

echo
echo '# A tampered object is refused on receipt: the hash IS the integrity check.'
printf 'X' | dd of="$SHARED/objects/$RH" bs=1 seek=3 count=1 conv=notrunc 2>/dev/null
echo '$ pp --transport-pull object <hash> shared' # after a byte flips in transit'
HOME="$B" "$PP" --transport-pull object "$RH" "$SHARED" 2>&1

rm -rf "$HOME"
```

```
$ pp node.pp # on machine A: the node body runs
[info] building greeter.o
42
node key: 2c46800e4d1eb254a42f78747d51201ff5579a023eb5c53d101be4fda5102b57
result hash: 8e5cc2a18337e30e25a6669b92076e5ce93431eefa209a68ddd13f2dd29fb36b
```

```
$ pp node.pp # on machine B, independently: SAME node key
[info] building greeter.o
42
machine B computed the same key: yes
```

```
# A publishes the result object by hash; B pulls it by that hash.
$ pp --transport-push object <hash> shared # on A
$ pp --transport-pull object <hash> shared # on B, re-hash-verified
byte-identical across the transport: yes
```

The two machines never agreed on a key. They derived the same one, because the key is a function of the code and nothing else. That is what makes a build computed “here” usable “there”. The result is addressed by what it is, not by where it was produced.

9.2 What is real, and what is stubbed

Be clear about the ground truth. The by-hash sync, the re-hash-on-receipt, the signed-token authority checks, and remote placement are all implemented and tested. But they run over a local-directory loopback transport, with two separate `~/.pp` stores standing in for two machines. The real network transport (ssh) is a stub: every operation raises a clear “not yet” error. The security model lives above the transport: content hashes, HMAC-signed capability tokens, authority gating. So a plaintext local-dir copy is deliberately the harder case to get right, and it is the case the tests exercise. A real cluster still needs the ssh transport underneath and a membership policy on top. The mechanisms they would carry are proven on the loopback.

9.3 Host-qualified identity

A running process is not a pure value. Here the “same identity everywhere” story inverts in exactly the right way: a process on host B is a different cell than the same process on host A. Host-qualified identity names world state, “the greeter running on B”, not a location you dispatch to. The desired-state map from the previous chapter generalizes one level for this, from `{domain → desired}` to `{host → {domain → desired}}`. `--member-name B` selects host B’s slice and hands it to the unchanged reconciler. A member is then just `pp --watch --supervise --member-name B` on its own slice: the single-machine supervisor, verbatim, fed a slice of a cluster-wide value. Naming a host absent from the map is a hard error, not a silent no-op.

Membership is ambient configuration: which machines exist, and where their stores are. It is read from `~/.pp/cluster/members` or `$PP_CLUSTER_MEMBERS`. It is deliberately not a capability. An address is not an authority ceiling, the same distinction the language draws everywhere between location and permission. Authority to act across the cluster travels separately, as a signed token minted from the same `--grant` grammar you use locally.

9.4 Trace merge, audit, and growth

Because a result carries its verifying trace, syncing a result syncs its evidence. A machine that pulls a node’s result also pulls the `(cell, hash)` observations that justify it. It re-verifies them against its own world before serving a hit. The authority checks survive the crossing. A hit served across the wire is gated on the requesting token’s authority over the trace’s read closure, exactly as a local hit is gated on the caller’s. A narrow caller cannot launder a broad read through a cached aggregator just because the bytes arrived over a network. And `pp why`, run on a machine that pulled a trace, redacts any cell the local caller is not authorized to see, identically to a purely local run. Audit redacts. It never lies.

A shared store only grows, so growth is bounded explicitly. `pp gc` keeps the last few successful reconcile passes as roots and sweeps everything their closures do not reach. It is never automatic, and never runs during a scheduled force. A trace does not record its children by key, so there is no on-disk graph to walk. Instead `gc` replays each root as an ordinary `pp` run. Every cache hit it makes marks its objects, traces, and blobs live. The sweep biases toward keeping: a root that fails to replay aborts the whole sweep. Over-retention is safe, and deleting live data is the only hazard.

10 The two back ends

`pp` runs a program in one of two ways. The tree-walker interprets the AST directly. It is small, obvious, and the reference for what every form means: the oracle. The bytecode VM compiles that same AST to a compact stack machine with $O(1)$ indexed locals and tail-call optimization, then runs the result. The

two engines share one reader, one set of value and hashing definitions, and one pool of runtime state. They are held to a single rule: on any program they must produce identical output. A disagreement between them is a bug, never a feature. The tree-walker says what the answer is, and the VM must match it.

That rule is what makes having two engines pay off. The tree-walker stays simple enough to trust. The VM can be optimized freely, because any divergence it introduces is mechanically caught against the oracle.

10.1 Running each engine

`pp file.pp` runs the tree-walker. `pp --bytecode file.pp` compiles the same file and runs it on the VM. Same source, same output:

be-bytecode.sh

```
#!/bin/sh
# The same program under each back end. `pp` uses the tree-walker; `pp
# --bytecode` compiles to the stack machine. Their output is identical.
export HOME=$(mktemp -d)

cat > "$HOME/prog.pp" <<'PP'
def fib(n) { if n < 2 { n } else { fib(n - 1) + fib(n - 2) } }

print(fib(15))
PP

echo '$ pp prog.pp'
"$PP" "$HOME/prog.pp"

echo '$ pp --bytecode prog.pp'
"$PP" --bytecode "$HOME/prog.pp"

rm -rf "$HOME"
```

```
$ pp prog.pp
610
$ pp --bytecode prog.pp
610
```

10.2 --diff: check them against each other

`pp --diff file.pp` runs the program under both back ends and compares the value every top-level form returned. It exits 0 when they agree. It exits 1 when they diverge, naming the file and printing both value lists. Because each engine executes the program, its own print output appears twice. The exit code is the verdict.

be-diff.sh

```
#!/bin/sh
# `pp --diff` runs a program under BOTH back ends and compares the values
# every top-level form returned. It exits 0 when they agree, 1 (naming the
# file and both value lists) when they diverge. The program's own output
# therefore appears twice – once per engine – and the exit code is the verdict.
export HOME=$(mktemp -d)

cat > "$HOME/prog.pp" <<'PP'
# a tail-recursive sum – both back ends must run it in constant stack
# and return the same value
def sum-to(n, acc) { if n = 0 { acc } else { sum-to(n - 1, acc + n) } }

print(sum-to(100000, 0))
PP

echo '$ pp --diff prog.pp'
"$PP" --diff "$HOME/prog.pp"
echo "exit: $?"

rm -rf "$HOME"
```

```
$ pp --diff prog.pp
5000050000
5000050000
exit: 0
```

The tail-recursive loop above also exercises a property both engines must share: tail calls run in constant stack. So the count to 100000 neither overflows nor differs between back ends.

Note the scope of the comparison. `--diff` compares the returned values of the top-level forms. The test suite (`dune runtest`) also diffs the two engines' stdout. So it also catches print- and effect-ordering divergences that leave the return values untouched. The two checks are complementary.

10.3 The differential fuzzer

`--diff` only checks the programs you hand it. To find divergences you would never think to write, `pp` ships a differential fuzzer (`tools/fuzz.ml`). It generates random programs, runs each under both back ends, and asserts identical observable behavior: same exit status, same stdout, same effect log. A mismatch is deduplicated by signature, shrunk to a minimal repro, and written to a failure directory. The generator is fully deterministic. Program `i` under a given seed is always the same program, so any finding reproduces exactly.

The fuzzer runs against two grammars. `core` covers the forms both back ends must always agree on: literals, `if/do/let/let*`, `fn` and application, `def`, the arithmetic and collection builtins, and the `stdlib` list functions. Any non-agreement on `core` is a real bug and fails CI. `full` adds the harder surface: type annotations, modules, effects and handlers, deep recursion, quotation, and `defmacro`. This is the project's single most valuable correctness asset. Every optimization the VM gains is only as trustworthy as the oracle it is diffed against. The fuzzer is what keeps that diff honest across the whole language, not just the examples someone remembered to test.

11 Language reference

This appendix gives the full detail the language chapter summarized. Each section covers one construct and shows it running. `pp` is a Lisp-1: functions and variables share one namespace, and evaluation is inner-first and strict — a form's arguments are values by the time it runs.

The special forms — the syntax the reader treats specially, rather than ordinary function calls — are: `if`, `do`, `let`, `let*`, `fn`, `def`, `defnode`, `quote`, `quasiquote` with `unquote` and `unquote-splicing`, `and`, `or`, `delay`, `force`, `node`, `perform`, `with-handler`, `with-caps`, `with-config`, `config`, `module`, `import`, `load`, `load-module`, `island`, and the `:` type annotation. `defmacro` is recognized structurally at the top level. Everything else is a function.

11.1 Value types

Integers, floats, strings, the booleans `true` and `false`, `nil`, keywords (`:name`), and quoted symbols (`quote { name }`) are the atoms. `print` shows strings quoted and prints one value per call.

ref-atoms.pp

```
# Integers, floats, strings, the booleans, nil, keywords, and quoted symbols.
# `print` shows strings quoted; everything else prints as it reads.
print(42)
print(3.14)
print("text")
print(true)
print(false)
print(nil)
print(:keyword)
print(quote { sym })
```

```
$ pp ref-atoms.pp
```

```
42
3.14
"text"
true
false
nil
:keyword
sym
```

The collections are lists, vectors, maps, and sets, each with its own printed form.

ref-collections.pp

```
# Lists, vectors, maps, and sets — each with its own reader syntax.
print(list(1, 2, 3))
print([1, 2, 3])
print({:a -> 1, :b -> 2})
print(hash-set(1, 2, 3))
```

```
$ pp ref-collections.pp
```

```
(1 2 3)
[1 2 3]
{:a 1, :b 2}
#{1 2 3}
```

11.2 Arithmetic and comparison

`+` and `*` are variadic. `-` and `/` take exactly two arguments. `mod` is the integer remainder.

ref-arith.pp

```
# `+` and `*` are variadic; `-` and `/` take exactly two arguments; `mod`  
# is the remainder.  
print(+(1, 2, 3))  
print(10 - 3)  
print(*(2, 3, 4))  
print(20 / 4)  
print(17 mod 5)  
  
$ pp ref-arith.pp
```

```
6  
7  
24  
5  
2
```

Comparisons are variadic and chain left to right — `<(1, 2, 3)` means `1 < 2` and `2 < 3`.

ref-compare.pp

```
# Comparisons are variadic and chain left to right: (< 1 2 3) is 1<2 and 2<3.  
print(2 = 2)  
print(<(1, 2, 3))  
print(>(3, 2, 1))  
print(<=(1, 1, 2))  
print(3 >= 3)  
  
$ pp ref-compare.pp
```

```
true  
true  
true  
true  
true
```

11.3 Booleans and conditionals

Only `nil` and `false` are falsy; every other value, including `0`, is truthy. `and` and `or` short-circuit and return the deciding value rather than a boolean.

ref-bool.pp

```
# Only nil and false are falsy; everything else (including 0) is truthy.
# `and`/`or` short-circuit and return the deciding value, not a boolean.
print(if 1 { if 2 { 3 } else { false } } else { false })
print(if false { true } else if nil { true } else { 5 })
print(if true { false } else { false })
print(not(nil))
print(not(0))

$ pp ref-bool.pp
```

```
3
5
false
true
false
```

if is an expression that evaluates exactly one branch. The untaken branch never runs — no effects, no errors.

ref-if.pp

```
# `if` is an expression and evaluates exactly one branch — the untaken one
# never runs, so the error here is never raised.
print(if 5 > 0 { "pos" } else { "neg" })
print(if false { error("never") } else { "safe" })

$ pp ref-if.pp
```

```
"pos"
"safe"
```

11.4 Bindings

let bindings are mutual: every binding is visible in every right-hand side and in the body, regardless of the order written.

ref-let.pp

```
# `let` bindings are mutual: every binding sees every other, whatever the
# textual order.
print(let (y = x + 1, x = 1) { y })

$ pp ref-let.pp
```

```
2
```

let* is the sequential form. Each binding sees only the ones above it, so a name may shadow itself.

ref-letstar.pp

```
# `let*` is sequential: each binding sees only the ones written above it,  
# so a name can shadow itself.  
print(let* (x = 1, x = x + 1) { x })  
  
$ pp ref-letstar.pp
```

```
2
```

At the top level, `let name = value` evaluates the value once and binds it; `def name(args...) { body }` defines a function.

ref-def.pp

```
# At the top level, (def name value) evaluates the value once and binds it;  
# (def (name args) body) defines a function.  
let answer = 6 * 7  
print(answer)  
def square(x) { x * x }  
print(square(7))  
  
$ pp ref-def.pp
```

```
42  
49
```

11.5 Functions

`fn` produces an anonymous function; `def` with a list head names one. There is no `loop` keyword — recursion is the loop, and both back ends run tail calls in constant stack.

ref-fn.pp

```
# `fn` makes an anonymous function; recursion is the only loop.  
print((fn(x) { x + 1 })(41))  
def fact(n) { if n = 0 { 1 } else { n * fact(n - 1) } }  
  
print(fact(5))  
  
$ pp ref-fn.pp
```

```
42  
120
```

11.6 Sequencing

`do` forces each form in order for its effects and returns the value of the last.

ref-do.pp

```
# `do` runs each form in order for its effects and returns the last value.
print(do {
  print("step 1")
  print("step 2")
  99
})

$ pp ref-do.pp
```

```
"step 1"
"step 2"
99
```

11.7 Type annotations

Annotations are optional gradual claims, written with `:`, and checked when the function's body runs — not before. A well-typed call passes through unchanged.

ref-types.pp

```
# Type annotations are optional and checked when the body runs. A well-typed
# call passes through unchanged.
def inc(n: int): int { n + 1 }
print(inc(41))

$ pp ref-types.pp
```

```
42
```

A mismatch names the offending value and its source location.

ref-type-error.pp

```
# A mismatch names the offending value and location.
def inc(n: int): int { n + 1 }
inc("oops")

$ pp ref-type-error.pp
```

```
pp: error: type mismatch: expected int, got "oops" at ref-type-error.pp:2
```

11.8 Quotation

Braces are pp's surface syntax; s-expressions are its AST. `quote { ... }` is the bridge: it turns the one form inside into that AST, as data. This is the same position Elixir takes — homoiconicity lives at the AST layer, not the surface. The brace text you type is not itself a data structure, but the tree it reads to is, and `quote` hands you that tree. Quotation is total: any form the reader accepts, `quote` turns into data.

Quasiquote builds structure with holes. The body of `quasiquote { ... }` is a template written in ordinary brace syntax, denoting the s-expression data it reads to; `unquote(e)` fills one hole with a computed value, `splice(e)` splices a list into a list position.

ref-quote.pp

```
# Braces are the surface; s-expressions are the AST. quote { ... } turns
# the form inside into that AST, as data. Quasiquote is a template:
# unquote(e) fills one hole, splice(e) splices a list into a list position.
print(quote { if a { b } else { c } })
print(quasiquote { f(1, unquote(1 + 1), splice(list(3, 4))) })

$ pp ref-quote.pp
```

```
(if a b c)
(f 1 2 3 4)
```

11.9 Laziness

delay suspends a computation; force runs it and memoizes the result, so the body runs at most once however many times it is forced. force is the identity on non-thunks.

ref-delay.pp

```
# `delay` suspends a computation; `force` runs it and memoizes the result,
# so the body runs at most once however many times it is forced.
let (t = delay(do {
  print("run once")
  42
})) {
  print(force(t))
  print(force(t))
}

$ pp ref-delay.pp
```

```
"run once"
42
42
```

11.10 Effects and config

perform dispatches an effect to the nearest handler installed by with-handler. Handling an effect yourself needs no capability.

ref-perform.pp

```
# `perform` dispatches to the nearest handler installed with `with-handler`.
print(with-handler(double = fn(x) { x * 2 }) { perform double(21) })

$ pp ref-perform.pp
```

```
42
```

Config is ambient, dynamically-scoped data, distinct from capabilities. config reads a key, with an optional default for when it is absent.

ref-config.pp

```
# Config is ambient, dynamically-scoped data. `config` reads a key, with an
# optional default when the key is absent.
print(with-config({:host -> "db1", :port -> 5432}) { config(:host) })
print(with-config({:host -> "db1"}) { config(:timeout, 30) })

$ pp ref-config.pp
```

```
"db1"
30
```

11.11 Macros

`defmacro` defines a macro: it receives its arguments as unevaluated forms — s-expression data, the same trees `quote` yields — and returns a new form, expanded before either back end sees it. You write the macro in braces; it consumes and produces the AST. A `quasiquote` `{ ... }` template is the usual way to assemble the expansion: ordinary brace syntax with `unquote(e)` holes where the caller's forms go.

ref-defmacro.pp

```
# `defmacro` receives its arguments as unevaluated forms (sexpr data) and
# returns a new form, expanded before either back end sees it. The body of
# quasiquote { ... } is a template in ordinary brace syntax.
defmacro unless(test, body) {
  quasiquote { if unquote(test) { nil } else { unquote(body) } } }
print(unless(false, "ran"))
print(unless(true, "ran"))

$ pp ref-defmacro.pp
```

```
"ran"
nil
```

Because a macro's real domain is the AST, the s-expression notation remains a first-class way to write one: a .ppl file is the same language in AST-native form — `pp` reads it with the s-expression reader, and the template `(if ,test nil ,body)` there builds the identical tree the brace template above does. Data operations work on either origin: where `quasiquote` gets awkward, assemble the form directly with `list/cons` — `list(quote { + }, x, x)` is a complete macro body. `gensym` supplies fresh names for any binding a macro introduces; `pp` macros are unhygienic, so the discipline is manual.

11.12 Modules

A module is a block whose definitions become a value. `import` merges that value's exports into the current scope; a module's own body is a fresh scope.

ref-module.pp

```
# A module is a block whose definitions become a value; `import` merges that
# value's exports into the current scope.
let (m = module { def double(x) { x * 2 }
  let tau = 6
}) {
  import(m)
  print(double(tau))
}

$ pp ref-module.pp
```

12

11.13 Lists

The core list operations are builtins and need no library: `list`, `cons`, `car`, `cdr`, and `map`.

ref-list-ops.pp

```
# List builtins need no library: list, cons, car, cdr, and map are primitive.
print(list(1, 2, 3))
print(cons(0, list(1, 2)))
print(car(list(1, 2, 3)))
print(cdr(list(1, 2, 3)))
print(map(fn(x) { x * x }, list(1, 2, 3, 4)))

$ pp ref-list-ops.pp
```

```
(1 2 3)
(0 1 2)
1
(2 3)
(1 4 9 16)
```

`stdlib/list.pp` adds the higher-order functions — `foldl`, `foldr`, `filter`, `range`, `take`, `drop`, `reverse`, `length`, `nth`, `append`, `member?`, and `each`. A program loads it with `load("stdlib/list.pp")`.

ref-list-stdlib.sh

```
#!/bin/sh
# stdlib/list.pp adds the higher-order list functions. It resolves relative to
# the repo root, so the script runs pp from there.
export HOME=$(mktemp -d)
cd "$(dirname "$(dirname "$PP")")"

"$PP" /dev/stdin <<'PP'
(load "stdlib/list.pp")
(print (range 1 6))
(print (foldl + 0 (range 1 6)))
(print (filter (fn (x) (> x 2)) (list 1 2 3 4)))
(print (take 2 (list 10 20 30 40)))
(print (drop 2 (list 10 20 30 40)))
(print (reverse (list 1 2 3)))
(print (length (list 1 2 3 4)))
(print (nth 1 (list 10 20 30)))
(print (append (list 1 2) (list 3 4)))
(print (member? 2 (list 1 2 3)))
PP

rm -rf "$HOME"
```

```
(1 2 3 4 5)
15
(3 4)
(10 20)
(30 40)
(3 2 1)
4
20
(1 2 3 4)
true
```

11.14 Maps

Maps are immutable. The builtins construct one (`hash-map`), look up a key (`hash-map-get`), and return new maps or their components (`map-insert`, `map-remove`, `map-keys`, `map-vals`).

ref-map-ops.pp

```
# Map builtins: construct, look up, insert, and read keys/values. Maps are
# immutable – map-insert and map-remove return new maps.
let m = {:a -> 1, :b -> 2}
print(hash-map-get(m, :a))
print(map-insert(m, :c, 3))
print(map-keys(m))
print(map-vals(m))
print(map-remove(m, :a))

$ pp ref-map-ops.pp
```

```
1
{:c 3, :a 1, :b 2}
(:a :b)
(1 2)
{:b 2}
```

stdlib/map.pp adds map-has? and map-merge on top of those builtins; it depends on stdlib/list.pp, so load that first.

ref-map-stdlib.sh

```
#!/bin/sh
# stdlib/map.pp builds on the map builtins and on stdlib/list.pp, so load that
# first.
export HOME=$(mktemp -d)
cd "$(dirname "$(dirname "$PP")")"

"$PP" /dev/stdin <<'PP'
(load "stdlib/list.pp")
(load "stdlib/map.pp")
(def m (hash-map :a 1 :b 2))
(print (map-has? m :a))
(print (map-has? m :z))
(print (map-merge m (hash-map :b 9 :c 3)))
PP

rm -rf "$HOME"
```

```
true
false
{:c 3, :b 9, :a 1}
```

11.15 Strings

The string builtins cover appending, length, splitting, substrings, index-of, trimming, and conversion to and from numbers.

ref-string-ops.pp

```
# String builtins: append, length, split, substring, index-of, trim, and
# number conversion.
print(string-append("foo", "bar"))
print(string-length("hello"))
print(string-split("a,b,c", ","))
print(string-sub("hello", 1, 3))
print(string-index("hello", "ll"))
print(string-trim(" hi "))
print(number->string(42))
print(string->number("3.14"))

$ pp ref-string-ops.pp
```

```
"foobar"
5
("a" "b" "c")
"ell"
2
"hi"
"42"
3.14
```

stdlib/string.pp adds string-join, starts-with?, ends-with?, and lines.

ref-string-stdlib.sh

```
#!/bin/sh
# stdlib/string.pp adds string helpers over the string builtins.
export HOME=$(mktemp -d)
cd "$(dirname "$(dirname "$PP")")"

"$PP" /dev/stdin <<'PP'
(load "stdlib/string.pp")
(print (string-join ", " (list "a" "b" "c")))
(print (starts-with? "hello" "he"))
(print (ends-with? "hello" "lo"))
(print (lines "a
b
c"))
PP

rm -rf "$HOME"
```

```
"a, b, c"
true
true
("a" "b" "c")
```

11.16 Type predicates

Each predicate forces its argument and tests its shape. The full set is `int?`, `float?`, `string?`, `bool?`, `keyword?`, `symbol?`, `pair?`, `vector?`, `map?`, `set?`, and `fn?`.

ref-predicates.pp

```
# Type predicates force their argument and test its shape.
(print(int?(3))
(print(string?("x"))
(print(keyword?(:k))
(print(pair?(list(1))
(print(map?({:a -> 1}))
(print(fn?(car))
(print(vector?([1]))
```

```
$ pp ref-predicates.pp
```

```
true
true
true
true
true
true
true
```

12 Command-line reference

This appendix lists every flag `pp` accepts, grouped by what you use it for. It comes from the argument parser in `src/main.ml`. Where the built-in `pp --help` and this table disagree, the source wins. A few flags marked internal are dispatch machinery that `pp` invokes on itself. They are here for completeness; you should not need to type them by hand.

Flags compose the way you would expect: `--bytecode`, `--grant`, `--schedule`, and `--watch` all layer onto whichever run mode you pick. Anything after a bare `--` becomes the program's own argument vector, which you read with `argv()`.

12.1 Running programs

Invocation	Meaning
<code>pp</code>	Start the REPL (tree-walker).
<code>pp <file.pp></code>	Read and run a source file; the program's value is its last top-level form.
<code>pp run <file></code>	Same as <code>pp <file></code> — <code>run</code> is an explicit subcommand spelling.
<code>pp -e '<expr>'</code>	Evaluate one expression string and print each top-level value. Not compatible with <code>--diff</code> .
<code>pp --once <file.pp></code>	Run once and exit. A no-op: this is already the default; the flag is for symmetry with <code>--watch</code> .
<code>pp -- <args...></code>	Everything after <code>--</code> becomes the program's <code>argv</code> , read via the <code>argv()</code> builtin.
<code>pp --version</code> , <code>pp -v</code>	Print the version and exit.
<code>pp --help</code> , <code>pp -h</code>	Print the usage summary and exit.

12.2 Back ends

`pp` has two interpreters over one shared store. The tree-walker is the default; `--bytecode` selects the VM. Both eliminate tail calls and share the same node cache, so a program's value is identical either way.

Flag	Meaning
<code>--bytecode <file.pp></code>	Run via the bytecode VM instead of the tree-walker.
<code>--diff <file.pp></code>	Run the file on both back ends and compare every top-level value. Implies <code>--bytecode</code> ; exits 1 on any mismatch, printing both sides. Not supported with <code>-e</code> .

12.3 Capabilities and grants

Side effects that touch the world require a capability, and capabilities enter a run only here, on the command line. `--grant` is repeatable; each grant is a colon-separated spec. Paths are canonicalized at the grant, so downstream authority checks and cell identities see one spelling.

Grant spec	Authority granted
<code>fs:<path>:ro</code>	Read access to a filesystem path (and everything under it).
<code>fs:<path>:rw</code>	Read and write access to a path.
<code>fs:<path>:wo</code>	Write-only access to a path.
<code>net:<host></code>	Network access to a host, any port.
<code>net:<host>:<port></code>	Network access to a host on one port.
<code>secret:<path></code>	Read a path as a sealed value (bytes never enter the store; <code>unseal</code> is the one way out).
<code>process</code>	Spawn, signal, and reap child processes.

Usage: `pp --grant fs:/tmp/build:rw --grant process <file.pp>`. Any of the world-touching primitives (`slurp`, the `fs/proc` domains, network effects) raises a capability error if the matching grant is absent.

12.4 The store and auditing

Every node's result is content-addressed and cached in `~/ .pp/store`. These flags inspect, bypass, audit, and reclaim that store.

Flag	Meaning
<code>why <file.pp> (--why)</code>	Run the file and explain each node's cache hit or miss. Capability-filtered: you only see nodes you had authority to observe.
<code>--no-cache <file.pp></code>	Skip cache reads and recompute every node; results are still written to the store.
<code>--check <file.pp></code>	Determinism audit: run each node twice and flag any whose result differs (a volatile node). With a non-serial <code>--schedule</code> , also re-runs the whole program serially and compares the desired-state hash. Exits 1 if anything is flagged.
<code>graph</code>	Print the cell→node dependency graph reconstructed from stored traces. Needs no file.
<code>gc</code>	Explicit store garbage collection (never automatic). Marks the store reachable from the last N reconcile/supervise epochs by replaying them, then sweeps the rest.
<code>gc --gc-keep-epochs <N></code>	Keep the last N epochs (a small built-in default; $N > 0$).
<code>gc --gc-grace-seconds <S></code>	Spare anything younger than S seconds from the sweep ($S \geq 0$).

12.5 Reconcile, watch, and supervise

These turn a program's value into managed state in the world. `--reconcile` materializes a `{relpath → content}` map as a file tree; `--supervise` drives a `{name → spec}` map of long-running processes. Both auto-load their domain policy from the `stdlib` and both are effectful, so both need grants.

Flag	Meaning
<code>--reconcile <root> <file.pp></code>	Materialize the program's map value as a file tree under <code><root></code> , creating/updating/deleting by content hash. Auto-loads <code>stdlib/domain-fs.pp</code> . Needs <code>fs:<root>:wo</code> (or <code>:rw</code>).
<code>--supervise <file.pp></code>	Converge the program's process map: start/restart/stop services to match it. Auto-loads <code>stdlib/domain-proc.pp</code> . Needs <code>process</code> . Pair with <code>--watch</code> to keep services alive.
<code>--watch <file.pp></code>	Run, then poll the observed cells and re-evaluate whenever one changes. Combines with <code>--reconcile</code> / <code>--supervise</code> to re-converge on drift.
<code>--stabilize</code>	With <code>--watch</code> : propagate changes by dirtying only the affected nodes (push stabilization) rather than re-running cold.
<code>--watch-interval <s></code>	Poll interval for <code>--watch</code> , in seconds (default 1.0).
<code>--fenced-policy retry abort ask</code>	How to treat a fenced (non-convergent) action left in an unknown state by a prior crash. Default abort.

12.6 Islands

Islands are content-addressed module pins — a git ref resolved to a hash and recorded inline in the source. These flags manage the pins.

Flag	Meaning
island-pins <file.pp>	List the file's island forms with their pin and cache status. Does not run the program.
--update <file.pp>	Re-resolve each island ref and rewrite the inline pins in place, then run. Implies --fetch-islands.
--fetch-islands	Allow a git fetch for island pins not already cached (off by default — an offline run only uses what the store already holds).

12.7 Scheduling and distribution

By default, pp computes node misses serially, in-process. --schedule changes where and how they run, up to placing them on other cluster members. The M5 cluster flags establish trust (cluster-init, --mint-token) and move desired state between machines by hash (--publish-object, --desired-object).

Flag	Meaning
--schedule serial	Compute node misses one at a time, in-process (the default).
--schedule parallel:<N>	Fan out independent misses across up to N workers.
--schedule race:<N>	Race up to N redundant computations, taking the first to finish.
--schedule remote:<member>	Place misses on a named cluster member.
cluster-init	Mint ~/.pp/cluster/{secret,id} — the cluster's trust anchor.
--mint-token <out> <ttl-secs>	Mint a signed cluster token into <out>, carrying whatever --grant specs accompany it, valid for <ttl-secs>.
--member-name <n> <file.pp>	Host-qualified distribution: treat the desired state as a {host → {domain → desired}} map and converge only host <n>'s slice. Combine with --reconcile/--supervise.
--publish-object <shared-root> <file.pp>	Run the program, store its fully-forced value (and every blob: ref it names) into a shared local-dir store by hash, and print the hash.
--desired-object <hash> <shared-root>	Pull a published value by <hash> from <shared-root> and converge it directly — never runs a program to derive it. Takes --member-name and the reconcile/supervise flags.

12.7.1 Internal transport (not for hand use)

pp invokes these on itself to move artifacts and dispatch remote work. They are low-level and explicit by design; a normal run never types them.

Flag	Meaning
--transport-push object blob trace <id> <root>	Copy one artifact into a shared local-dir store.

Flag	Meaning
--transport-pull object blob trace <id> <root>	Copy one artifact out of a shared local-dir store (re-hash-verified on ingest).
--serve-hit <key> <token-file> <shared-root> <reply-file>	Capability-gated cache lookup served to a dispatcher.
--recv-hit <reply-file> <shared-root>	Ingest a serve-hit reply.
--remote-node <token> <pins> <root> <keys> <reply>	The cluster-member side of remote placement; authority comes from the verified token, not --grant.
--gc-mark <outfile>	Used only by pp gc's own replay subprocess to record the live set.

12.8 Observation pinning

These capture and replay the exact set of world observations a run made — the seam that lets one machine pin another's reads (and the basis of remote placement's pin wire).

Flag	Meaning
--pin-file <path> <file.pp>	Preseed the run's cell observations and probe values from a file of (pin ...) / (pin-probe ...) lines before the program runs, so it never observes its own disk for a pinned cell.
--dump-pins <path> <file.pp>	After the run, write every observed cell and probe value to <path> as (pin ...) / (pin-probe ...) lines. A probe value that is not plain data (code, a handle, a sealed secret) is skipped with a warning.

The pin-file's (pin ...) / (pin-probe ...) lines are their own small wire format (src/remote.ml's parse_pin_line), not pp source — they are not read by either pp reader, so they keep their fixed parenthesized shape regardless of the surface a program is written in.

13 Builtins and the standard library

This appendix indexes the two layers of pp's vocabulary that are not special forms: the builtins compiled into the binary (src/primitives.ml) and the standard library written in pp itself (stdlib/*.pp). Special forms — if, let, def, fn, delay, force, node, perform, handle, quote, and the rest — belong to the language reference, not here.

Signatures use pp calling syntax: name(arg, ...). A trailing ... marks a variadic position; [arg] marks an optional one. Unless noted, a builtin forces the arguments it inspects; cons, list, hash-map, and map are deliberately lazy in the values they carry.

13.1 Core builtins

These are always in scope — no load needed.

13.1.1 Arithmetic and comparison

Signature	Description
+(n, ...)	Variadic sum; identity 0. Mixing ints and floats promotes to float.

Signature	Description
- (a, b)	Subtraction of exactly two numbers.
*(n, ...)	Variadic product; identity 1.
/(a, b)	Division of two numbers; the divisor must be non-zero.
mod(a, b)	Integer remainder; the divisor must be non-zero.
=(a, ...)	Structural equality; true iff all arguments are equal.
<(a, ...), >(a, ...)	Variadic chained ordering; true iff each adjacent pair is strictly ordered.
<=(a, ...), >=(a, ...)	As above, non-strict.

13.1.2 Lists and pairs

Signature	Description
cons(a, b)	Build a pair, storing both parts unforced (lazy).
car(p)	First element of a pair; nil for nil.
cdr(p)	Rest of a pair; nil for nil.
list(x, ...)	Build a proper list from its arguments (lazy in the elements).
map(f, lst)	Apply f to each element, consing the results without forcing them — the parallel fan-out point the scheduler batches on.
nil?(x)	True if x is nil.

13.1.3 Vectors, maps, and sets

Signature	Description
vector(x, ...)	Build a vector (lazy in the elements).
vector-get(v, i)	Element i of vector v; out-of-bounds is an error.
hash-map(k, v, ...)	Build a map from alternating keys and values; keys are forced, values stay lazy.
hash-map-get(m, k)	Value bound to k in m, or nil.
map-insert(m, k, v)	A new map with k bound to v (replacing any existing binding).
map-remove(m, k)	A new map without key k.
map-keys(m)	The keys of m as a list.
map-vals(m)	The values of m as a list.
hash-set(x, ...)	Build a set (lazy in the elements).

13.1.4 Type predicates

Each takes one argument, forces it, and returns a boolean.

Predicate	True when the value is...
int?(x), float?(x)	an integer / a float.
string?(x), bool?(x)	a string / a boolean.
keyword?(x), symbol?(x)	a keyword / a symbol.
pair?(x)	a pair or nil.
vector?(x), map?(x), set?(x)	a vector / a map / a set.
fn?(x)	a closure or builtin.
thunk?(x)	an unforced thunk (does not force its argument).
capability?(x)	a capability.

13.1.5 Strings and numbers

Signature	Description
<code>string-append(s, ...)</code>	Concatenate; non-string arguments are rendered first.
<code>string-length(s)</code>	Length of a string in bytes.
<code>string-split(s, sep)</code>	Split <code>s</code> on the single-character <code>sep</code> , dropping empty fields.
<code>string-index(s, sub)</code>	Index of the first occurrence of <code>sub</code> , or <code>nil</code> .
<code>string-trim(s)</code>	<code>s</code> with leading and trailing whitespace removed.
<code>string-sub(s, start, len)</code>	The <code>len</code> -byte substring at <code>start</code> ; out-of-bounds is an error.
<code>number->string(n)</code>	Decimal rendering of an int or float.
<code>string->number(s)</code>	Parse to an int, else a float, else <code>nil</code> .

13.1.6 I/O and control

Signature	Description
<code>print(x, ...)</code>	Deep-force each argument, print them concatenated with a trailing newline, and return <code>nil</code> .
<code>not(x)</code>	Logical negation; <code>nil</code> counts as false.
<code>error(msg)</code>	Raise an error with string message <code>msg</code> .
<code>exit([n])</code>	Terminate the run with status <code>n</code> (default 0).
<code>slurp(path)</code>	Read a file to a string. Needs an <code>fs read</code> grant (or a <code>secret</code> grant, which yields a sealed value); capability-free inside a node's scratch sandbox.
<code>argv()</code>	The program arguments after <code>--</code> , as a list of strings. Recorded as an <code>argv</code> : observation.
<code>env-get(name)</code>	The environment variable name, or <code>nil</code> . Recorded as an <code>env</code> : observation.

13.1.7 Filesystem observation

Capability-gated presence checks, recorded as `stat`: cells so a node recomputes exactly when the path appears, disappears, or changes kind — never reading contents.

Signature	Description
<code>file-exists?(path)</code>	True if anything exists at <code>path</code> . Needs an <code>fs read</code> grant.
<code>dir?(path)</code>	True if <code>path</code> is a directory. Needs an <code>fs read</code> grant.

13.1.8 Hashing and the content store

Signature	Description
<code>hash-value(v)</code>	A canonical structural content hash of any value — order-independent for maps and sets.
<code>hash-string(s)</code>	The SHA-256 hex digest of a string's raw bytes (pure; no store I/O).
<code>blob(s)</code>	Ingest bytes into the content store, returning a <code>"blob:<sha256>"</code> reference.
<code>blob-get(ref)</code>	The inverse of <code>blob</code> : the stored bytes for a <code>"blob:<hash>"</code> reference.

13.1.9 Capabilities

Signature	Description
<code>current-capabilities()</code>	Reify the ambient capability set as of the call — an observation of the ceiling, never a mint.
<code>cap-compose(c, ...)</code>	Compose several capabilities into one.

Signature	Description
<code>cap-restrict(cap, scope, [:ro :rw :wo])</code>	Narrow cap to a scope string and, optionally, a weaker mode.
<code>cap-none()</code>	The empty capability (grants nothing).

13.1.10 Effects, domains, and metaprogramming

Signature	Description
<code>register-domain({...})</code>	Register a write-domain from a spec map (<code>:name, :namespace, :o</code> cap). Script-tier only — not inside a node body.
<code>register-probe(name, observe-fn, read-cap)</code>	Register a read-only probe: a domain with no write authority. Script-tier only.
<code>probe(name)</code>	Read a registered probe's value, pinned once per pass and recorded as a <code>probe: cell</code> .
<code>fenced(kind, spec)</code>	Register a non-convergent (fenced) action for the reconciler to drain after convergent state is applied.
<code>unseal(v)</code>	Convert a sealed value to a string — the one sanctioned exit from <code>secret: bytes</code> .
<code>eval-pp(code)</code>	Parse and evaluate a code string in the current run's environment and macro table.
<code>apply-pp(fn, args)</code>	Apply <code>fn</code> to a list of already-evaluated arguments.
<code>force-deep(v)</code>	Fully force a value and everything reachable from it.
<code>read-string(src)</code>	Parse a source string to a quoted value (or a vector of them).
<code>gensym([prefix])</code>	A fresh, unwritable symbol for hygienic macro bindings.

The reader forms `quasiquote`, `unquote`, and `unquote-splicing` are also registered as builtins, and you normally reach them through the reader's `quasiquote` syntax; outside a `quasiquote` the latter two are an error. The `ppc-run` / `ppc-finish` family are self-hosting compiler scaffolding and are mostly unimplemented today.

13.2 Lists

`stdlib/list.pp` — load with `load("list.pp")`. Note that `map` is not here: it is a builtin, on purpose, so a loaded list library cannot shadow the scheduler's fan-out point.

Signature	Description
<code>filter(pred, lst)</code>	A lazy list of the elements satisfying <code>pred</code> .
<code>foldl(f, acc, lst)</code>	Left fold, strict in the accumulator.
<code>foldr(f, acc, lst)</code>	Right fold, lazy.
<code>range(start, end)</code>	The integers from <code>start</code> (inclusive) to <code>end</code> (exclusive).
<code>take(n, lst)</code>	The first <code>n</code> elements.
<code>drop(n, lst)</code>	<code>lst</code> without its first <code>n</code> elements.
<code>nth(n, lst)</code>	Zero-based element access; <code>nil</code> past the end.
<code>length(lst)</code>	Element count (strict — forces the whole list).
<code>each(f, lst)</code>	Apply <code>f</code> to each element for effect; returns <code>nil</code> .

Signature	Description
<code>append(a, b)</code>	Concatenate two lists (lazy in b).
<code>reverse(lst)</code>	Strict reversal.
<code>member?(x, lst)</code>	Structural membership test.

13.3 Maps

`stdlib/map.pp` — load with `load("map.pp");` depends on `list.pp`.

Signature	Description
<code>map-has?(m, k)</code>	Whether <code>k</code> is a key of <code>m</code> .
<code>map-merge(a, b)</code>	<code>a</code> with every binding of <code>b</code> inserted; <code>b</code> wins on collision.

13.4 Strings

`stdlib/string.pp` — load with `load("string.pp");` no dependencies.

Signature	Description
<code>string-join(sep, lst)</code>	Join a list of strings with <code>sep</code> between elements.
<code>starts-with?(s, prefix)</code>	Whether <code>s</code> begins with <code>prefix</code> .
<code>ends-with?(s, suffix)</code>	Whether <code>s</code> ends with <code>suffix</code> .
<code>lines(s)</code>	Split <code>s</code> into lines, dropping empty fields.

13.5 Domains

The domain policies are pp libraries that pp auto-loads for you: `stdlib/domain-fs.pp` under `--reconcile`, `stdlib/domain-proc.pp` under `--supervise` (each after `list.pp`, `map.pp`, and `string.pp`). The trusted mechanics they call — `tree-observe`, `materialize-file`, `proc-spawn`, and so on — are OCaml primitives reached only through `perform`. You normally interact with a domain through the registration entry point; the rest are its internal policy, listed here for readers of the source.

The one entry point you call directly:

Signature	Description
<code>register-fs-domain(root, write-cap)</code>	Register the filesystem domain rooted at <code>root</code> , converging a <code>{relpath → content}</code> map (content: an inline string or a blob: reference).
<code>register-proc-domain(write-cap)</code>	Register the process domain, converging a <code>{name → spec}</code> map (spec: <code>cmd/args/env/cwd</code>).

13.5.1 Filesystem policy internals (`domain-fs.pp`)

Signature	Description
<code>fs-blob-ref?(c)</code>	Whether content <code>c</code> is a blob: reference.
<code>fs-blob-ref-hash(c)</code>	The hash inside a blob: reference (dropping any <code>:x</code> suffix).
<code>fs-blob-ref-executable?(c)</code>	Whether a blob: reference carries the executable-bit <code>:x</code> suffix.
<code>fs-content-hash(c)</code>	A content's identity hash: the ref's hash, or hash-string of inline bytes.
<code>fs-content-bytes(c)</code>	The bytes for a content: pulled from the store for a blob: ref, else inline.
<code>fs-validate-rel-part(rel, part)</code>	Reject <code>..</code> in a desired path component.

Signature	Description
<code>fs-validate-rel(rel)</code>	Validate that a desired path is relative and traversal-free.
<code>fs-plan-item(kind, rel, content)</code>	Build one create/update/delete plan item.
<code>fs-diff-for(root)</code>	Return the pure (observed desired) → plan diff, closed over root.
<code>fs-apply-item(root, item)</code>	Materialize or remove one file per plan item.
<code>fs-apply-for(root)</code>	Return the apply function that runs each plan item.

13.5.2 Process policy internals (`domain-proc.pp`)

Signature	Description
<code>proc-state-key(name)</code>	The domain-state key under which a service's record is stored.
<code>proc-known-names()</code>	The domain's own index of tracked service names.
<code>proc-remember!(name, pid, spec, known)</code>	Record a running service and add it to the known set.
<code>proc-forget!(name, known)</code>	Clear a service's record and drop it from the known set.
<code>proc-observe-one(name)</code>	Observe one service: its spec if alive, <code>:stopped</code> if tracked-but-dead, else <code>nil</code> .
<code>proc-observe()</code>	Reap zombies, then observe every tracked service.
<code>proc-plan-item(kind, name, spec)</code>	Build one start/restart/stop plan item.
<code>proc-spec-eq?(a, b)</code>	Compare two specs by canonical hash-value (immune to map-order differences).
<code>proc-diff(observed, desired)</code>	The pure diff: start / restart / stop by spec change.
<code>proc-stop-current!(name)</code>	Stop a service at its currently-recorded pid.
<code>proc-apply-item(known, item)</code>	Apply one plan item, returning the updated known set.
<code>proc-apply(plan)</code>	Fold every plan item, then persist the known set once.